

Housing Price Analysis using Random Forest

Bangda Sun

May 14, 2017

We see that linear regression model works well on predicting house price, the final model's R^2 we get is approximately 0.874, which is very good in practice.

Next we try to move beyond linearity, one of the most important family of models in machine learning is tree models, including decision trees, bagging and random forest etc.

The predictors we used in the linear regression model includes:

sqftliving, bedrooms, bathrooms, grade, floors, waterfront, yrbuilt, lat, long, view, zipcode, condition.

1. Decision Tree

First we start from single decision tree,

```
library(tree)
tree1 = tree(price ~. -date-zipcode, data = training)
summary(tree1)

##
## Regression tree:
## tree(formula = price ~ . - date - zipcode, data = training)
## Variables actually used in tree construction:
## [1] "grade"      "lat"        "sqft_living" "waterfront" "yr_built"
## [6] "long"
## Number of terminal nodes: 12
## Residual mean deviance: 4.304e+10 = 3.046e+14 / 7076
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -4065000 -99520  -25680      0    68920  2195000
# Visualize the tree
plot(tree1)
text(tree1)
```




```
##      29) long > -122.179 34 9.649e+12 1426000 *
##      15) sqft_living > 7370 6 2.232e+13 4865000 *
```

Then we use cross validation to see whether pruning the tree will improve performance,

```
cv_tree1 = cv.tree(tree1)
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 3.3.2
```

```
ggplot() +
  geom_line(mapping = aes(x = cv_tree1$size, y = cv_tree1$dev)) +
  geom_point(mapping = aes(x = cv_tree1$size, y = cv_tree1$dev), size = 3) +
  labs(x = "size of tree", y = "dev")
```




We can see that the dev of tree will get minimum when the tree size is 12. Therefore, the performance of tree doesn't improve much if we prune the tree.

Next we calculate the test error of this tree,


```
predOfTree1 = predict(tree1, newdata = testing)
rmse1 = sqrt(mean((predOfTree1 - testing$price)^2))
rmse1
```

```
## [1] 224316.3
```

```
ggplot() +
  geom_point(mapping = aes(x = predOfTree1, y = testing$price), alpha = .2, size = 2) +
  geom_abline(slope = 1, intercept = 0) +
  labs(x = "prediction", y = "actual value")
```



```
ggplot() +
  geom_line(mapping = aes(x = 1:length(predOfTree1), y = predOfTree1 - testing$price)) +
  geom_abline(slope = 0, intercept = 0) +
  geom_abline(slope = 0, intercept = rmse1, linetype = "dashed", color = "red") +
  geom_abline(slope = 0, intercept = -rmse1, linetype = "dashed", color = "red") +
  labs(x = "observation", y = "price")
```



From this plot, we can see that many predictions are far away from their actual value. The error is still high. We need to seek models to predict more precisely.

2. Random Forest

We use ensemble methods to improve the performance of the tree model. One of the popular methods is random forest.

```
library(randomForest)

## randomForest 4.6-12
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:ggplot2':
##
##     margin
set.seed(1)
randForest = randomForest(I(log(price)) ~. -id-zipcode-date, data = training, mtry = 5,
                           importance = TRUE)
randForest


##
## Call:
```

```
## randomForest(formula = I(log(price)) ~ . - id - zipcode - date,      data = training, mtry = 5, imp
##           Type of random forest: regression
##           Number of trees: 500
## No. of variables tried at each split: 5
##
##           Mean of squared residuals: 0.03586097
##           % Var explained: 87.29


predOfrandForest = predict(randForest, newdata = testing)
rmse2 = sqrt(mean((testing$price - exp(predOfrandForest))^2))
rmse2
```

```
## [1] 142801.7
```

```
ggplot() +
  geom_point(mapping = aes(x = exp(predOfrandForest), y = testing$price), alpha = .2, size = 2) +
  geom_abline(slope = 1, intercept = 0) +
  labs(x = "prediction", y = "actual value")
```



```
ggplot() +
  geom_line(mapping = aes(x = 1:length(predOfrandForest), y = exp(predOfrandForest) - testing$price)) +
  geom_abline(slope = 0, intercept = 0) +
  geom_abline(slope = 0, intercept = rmse2, linetype = "dashed", color = "red") +
  geom_abline(slope = 0, intercept = -rmse2, linetype = "dashed", color = "red") +
  labs(x = "observation", y = "price")
```



```
importance(randForest)
```

##		%IncMSE	IncNodePurity
##	bedrooms	16.98279	17.682193
##	bathrooms	22.14704	71.858680
##	sqft_living	43.12953	325.628456
##	sqft_lot	48.61245	46.454551
##	floors	17.32741	14.573919
##	waterfront	25.23851	11.689565
##	view	34.66992	35.028867
##	condition	28.71623	14.200844
##	grade	33.68949	332.797907
##	sqft_above	26.46916	125.406629
##	sqft_basement	22.85204	26.940917
##	yr_built	47.72134	62.636548
##	yr_renovated	11.61159	3.968069
##	lat	182.13863	579.685747
##	long	77.01030	76.154332
##	sqft_living15	43.72745	171.641281
##	sqft_lot15	43.87485	53.835669

```
varImpPlot(randForest)
```

randForest

